

Question 1

Approach. The methodology behind this algorithm is to use a modified version of binary tree searching (*from CLRS, page 290*) to find the first node with the given prefix x and use a modified version of an in-order traversal (*from CLRS, page 288*) to traverse until the node keys no longer contain prefix x .

Algorithm 1: Prefix Search

```

1 FIND-STRINGS-WITH-PREFIX( $r, x$ ):
2    $f = \text{PREFIX-TREE-SEARCH}(r, x)$  // First node that matches prefix.
3   PREFIX-BASED-INORDER-WALK( $f, x$ )
4   return

5 PREFIX-BASED-INORDER-WALK( $r, x$ ):
6   if  $r \neq \text{NULL}$  and  $x == r.\text{key}[0 : \text{len}(x)]$  then
7     PREFIX-BASED-INORDER-WALK( $r.\text{left}, x$ )
8     print  $r.\text{key}$ 
9     PREFIX-BASED-INORDER-WALK( $r.\text{right}, x$ )
10  return

11 PREFIX-TREE-SEARCH( $r, x$ ):
12  if  $r == \text{NULL}$  or  $x == r.\text{key}[0 : \text{len}(x)]$  then
13    return  $r$ 
14  else if  $x < r.\text{key}[0 : \text{len}(x)]$  then
15    return PREFIX-TREE-SEARCH( $r.\text{left}, x$ )
16  else
17    return PREFIX-TREE-SEARCH( $r.\text{right}, x$ )

```

We assume in the above that the $<$ and $>$ operations are defined correctly for strings. (“ABC” $<$ “ACB”)

Proof of correctness and termination. The binary tree in-order traversal procedure is equivalent to that in CLRS except for the condition that the key of the current node, r , must contain the prefix x . It operates the same as described in CLRS except for when the key of r does not contain the prefix x , at which point it terminates on line 10.

The tree searching procedure is equivalent to that in CLRS except for checking whether the key of the current node, r , contains the given prefix, x , instead of matching a given key. Therefore it operates the same as described in CLRS.

Based on the correctness of the in-order traversal procedure and the tree searching procedure we can prove the correctness of the main algorithm by contradiction. *Assume by contradiction* that there is a node matching the prefix that this algorithm does not find, there are two cases:

1. *The node is not in the sub-tree of node f found on line 2.* This is not possible in a binary tree, since, if a prefix-matching node existed outside the sub-tree then it would have been selected as f by the tree searching procedure.
2. *The in-order traversal ends before reaching a prefix-matching node.* This is not possible as the result of an in-order traversal is sorted and in a sorted list of strings, those sharing a matching prefix will always be consecutively next to one another.

Time and space complexity. The time complexity of this algorithm for n strings where k contain prefix x is the sum of each of the procedures: the search algorithm is $O(\log(n))$ and the in-order traversal is $O(k)$ (since it will always stop running after k nodes). Therefore the time complexity overall is $O(\log(n) + k)$. The space complexity is $O(n)$ for the binary tree, since the algorithm does not store anything.

Question 2

Part A

Let $T(n)$ be the number of nodes in a full binary tree of height n (where $n \geq 1$). Then the recurrence representing $T(n)$ can be stated as:

$$T(n) = 2 \cdot T(n-1) + 1, \quad T(1) = 1$$

This is logical given that adding 1 level of height to a full binary tree consists of doubling the previous level (two children per node) and adding 1 for the root node in the tree.

Part B

We want to show that there are more leaf nodes than non-leaf nodes in a full binary tree. A leaf node is a node with no children, which in this case is the last level of the binary tree.

For a tree of height n , the number of non-leaf nodes is represented by $T(n-1)$, therefore the number of leaf nodes can be written as:

$$\begin{aligned} \# \text{ leafs} &= \# \text{ nodes} - \# \text{ non-leafs} \\ &= T(n) - T(n-1) \\ &= 2T(n-1) + 1 - T(n-1) \\ &= T(n-1) + 1 \end{aligned}$$

Since $T(n-1) + 1 > T(n-1)$, it is clear that the number of leaf nodes is always greater than the number of non-leaf nodes.

Question 3

1) $1 \bmod 11 = 1$

2) $22 \bmod 11 = 0$

3) $54 \bmod 11 = 10$

4) $13 \bmod 11 = 2$

5) $12 \bmod 11 = 1 \rightarrow \text{Collision!}$

(a) $(1+3(12) \bmod 4+1) \bmod 11 = 3$

6) $3 \bmod 11 = 3 \rightarrow \text{Collision!}$

(a) $(3+3(3) \bmod 4+1) \bmod 11 = 5$

7) $27 \bmod 11 = 5 \rightarrow \text{Collision!}$

(a) $(5+3(27) \bmod 4+1) \bmod 11 = 7$

Index	Data	Step #
0	22	2
1	1	1
2	13	4
3	12	5
4		
5	3	6
6		
7	27	7
8		
9		
10	54	3

Question 4

The greedy algorithm works as follows:

1. Sort set S by finishing time in ascending order
2. Take the first element in S , (s_1, f_1) , and determine which other elements overlap with this interval
3. Using greedy, determine time t within (s_1, f_1) s.t. all overlapping tournaments will be visited at t
4. Remove all intervals that overlap with (s_1, f_1) (including (s_1, f_1)) from the set
5. Resort the set
6. Repeat until the S is empty

Algorithm 2: Tournament Greedy Function

```

1  TOURNAMENT-GREEDY( $S$ ):
2     $ResultArray[n]$     // array of best times to go to the tournament
3    while  $S$  is not empty do
4       $SortByFinishTime(S)$ 
5       $t = S_1$ 
6       $i = 0$ 
7      while  $i \neq length(S)$  do
8        if  $s_i \leq f_1$  then
9          if  $s_i > t$  then
10              $t = s_i$ 
11              $DeleteFromSet(s_i, f_i)$ 
12          else
13              $i++$ 
14              $AppendToResultArray(t)$ 
15        end
16      end

```

Existence of Subproblems. After t is added to $ResultArray$ and all overlapping intervals (let's say k intervals) are deleted, the new set will be smaller, as $n < n - k$.

Proof of Correctness. Our goal is to minimize of times visited while visiting all siblings. Let R be the result array of the greedy solution.

Contradiction: Assume that R is not the optimal solution and instead there exists the optimal solution A where $length(A) < length(ResultArray)$.

$$\begin{array}{ccccccc} A_1 & A_2 & A_3 & \dots & & & \\ R_1 & R_2 & R_3 & \dots & R_i & \dots & R_k \end{array}$$

In order to satisfy the constraint of visiting all siblings, A_1 must lie in the interval (s_1, f_1) (when S is sorted in order by finishing time) as well as all other intervals that overlap with (s_1, f_1) . A_2 lies in the interval (s_m, f_m) where (s_m, f_m) is the minimum of $S - [(s_i, f_i) \supset A_1]$, where $i \in [0, n]$. Similarly, A_3 lies in the interval of the $\min(S - [(s_i, f_i) \supset [A_1, A_2]])$.

However, the greedy algorithm follows the same logic and also incorporates $R_1 \in (s_1, f_1)$, $R_2 \in \min(S - [(s_i, f_i) \supset R_1])$, etc. Since the greedy algorithm will always make the most locally optimal choice, it contradicts that there exists a more optimal algorithm and that $\text{length}(A) < \text{length}(\text{result})$.

Proof of Optimal Substructure. When the optimal solution is chosen for the earliest finish time, this will maximize the size of the number of k games removed from the set - in other words, this minimizes $n - k$ in the following subproblem. By choosing the optimal solution for each subproblem, we also minimize the number of potential subproblems, thereby minimizing our total number of trips to the tournament.


Question 5

The problem to solve is as follows: We want to travel from points x_1 to x_n on a trail and we can choose to stop at any point x_i , $i \in [2, n - 1]$ on the way. Every stop, x_i has an associated cost of stopping, b_i , and every two stops x_i and x_j have an associated cost of traveling without stopping, c_{ij} .

Approach. The larger problem from $[1, n]$ can be divided into small sub-problems where we are currently located at a stop x_i and must decide whether to continue without stopping to x_j . The goal is to choose the minimum cost between c_{ij} (not stopping) and the sum of b_i (stopping) with the cost of the rest of the journey. This can be expressed by the recursion below, let $m(i, j)$ represent the minimum cost of traveling between x_i and x_j .

$$m(i, j) = \begin{cases} 0, & \text{for } i=j \\ \min\{c_{ij}, b_i + m(i+1, j)\}, & \text{otherwise} \end{cases}$$

A recursive algorithm from the recurrence above would solve the problem however in order to reduce the number of computations we would store computed $m(i, j)$ values in an upper half table and store a 1 in location (i, j) to indicate a stop at i on the way to j . An example of these tables is shown below for $n = 3$, note that the path from $(1, 3)$ to $(3, 3)$ shows the stops that need to be made, which would be returned by the algorithm.



Arguments of correctness. Assume by contradiction that there is a cost smaller than the one found by the recursion, there are two cases:

1. *The algorithm stopped at some x_i when it should have continued to an x_j .* This is impossible since the algorithm always picks the minimum of c_{ij} and $b_i + m(i+1, j)$, so if there was a smaller c_{ij} it would have been chosen.
2. *The algorithm did not stop at an x_i when it should have.* By the same logic as above, if the algorithm always picks the minimum, if there was a b_i which resulted in a lower cost it would have been chosen.

Time complexity analysis. Since we are storing intermediate results in the table, the run-time of the algorithm is the time to fill the table, which is $O(n^2)$.